

Event driven simulation of a 2-dimensional gas

Sondre Duna Lundemo[†]

Department of Physics, Norwegian University of Science and Technology, Trondheim
Norway

TFY4235 - Computational physics
(*Last updated on February 15, 2021*)

An event driven approach is used to simulate a gas of hard discs in two dimensions. The framework built is used to study statistical properties of the gas at and towards equilibrium. The numerical results are compared with well known results from classical statistical mechanics.

Contents

1	Introduction and background	2
2	Overview of code	2
3	Results and discussion	5
3.1	Speed distribution at equilibrium	5
3.2	Mixture of two gases	6
3.3	Mixing in the presence of energy dissipation	6
3.4	Crater formation from projectile impact	7

1 Introduction and background

An event driven approach is used to simulate a gas of hard discs in two dimensions. The framework built is used to study statistical properties of the gas at and towards equilibrium. The numerical results are compared with well known results from classical statistical mechanics.

The central algorithm for event-driven simulations is the following [1]:

```

Set velocities and positions of all particles in the gas;
Choose a stop criterion;
for each particle in gas do
    Calculate if and when the particle will collide with all the other particles and
    the walls;
    Store all the collision times;
end
while not reached stop criterion do
    Identify the earliest collision;
    if collision is valid then
        Move all particles in straight lines until the earlieset collision;
        for each particle involved in collision do
            Calculate if and when the particle will collide with all the other
            particles and the walls;
            Store all the collision times;
        end
    else
        Discard collsion;
    end
end

```

Algorithm 1: Event driven simulation of a gas.

In the above algorithm, a collision is *valid* if the particle(s) involved in the collision has *not* collided since the time the collision time was stored.

2 Overview of code

In this section I briefly present how the algorithm is implemented and which considerations are put into the choices of data-structures. The code is written in `python`.

The main machinery of the code is collected in a class called `Gas` in `events.py`, which in essence is a collection of `N particles` and various methods to manipulate the coordinates of each particle to simulate the gas' time evolution according to the algorithm in section 1. The particles are represented by their coordinates in the extended configuration space by their coordinates in the extended configuration space, that is their position and velocity. Although the main algorithm presented in the introduction 1 invites to a fully object-oriented approach, I have chosen to restrain myself somewhat in that respect. For instance, I have chosen not to create an object representing a particle. The most prominent caveat preventing me from doing this is that it might affect the speed of the

calculations. By choosing not to create separate object representing each particle, I found it very simple to move the particles at each time step and also accessing them by ordinary slicing and indexing of arrays. Although it is possible to overload operators such that an array of self defined object can be added *like* `numpy`-arrays, I found this to be impractically slow. A quick test of adding the self-made particles compared to simply adding $4 \times N$ arrays establishes this observation quite firmly. The listings below shows the time spent on adding two arrays of 50 000 particles with each of the mentioned methods.

```
1 >>>%time particles_array_1 += particles_array_2
2 CPU times: user 443 us, sys: 101 us, total: 544 us
3 Wall time: 306 us
```

```
1 >>>%time particles_class_1 += particles_class_2
2 CPU times: user 34.1 ms, sys: 87 us, total: 34.2 ms
3 Wall time: 33.2 ms
```

The issue of finding the earliest event for each timestep is solved using a priority queue from the library `heapq` in `python`. This data structure allows for sorting any objects as long as they can be compared by the "less than"-operator.

The code itself is well documented and should be easy to understand on its own. However, I would like to point out some solutions that I found to work particularly well. The part of the code that undoubtedly is the most computationally heavy is the one devoted to calculating if and when the particles will collide with all others. The naive approach would be to iterate over each particle and do the calculation for each of them separately. In `python`, these kind of nested loops will often become impractically slow. I found considerable improvements through vectorising this calculation.

By essentially replacing the piece of calculation in listing 1 by that in 2 I was able to reduce one of the loops over all of the particles.

```
1 for j in range(self.N):
2     delta_x = self.particles[:2,j] - self.particles[:2,i]
3     delta_v = self.particles[2:,j] - self.particles[2:,i]
4     R_ij     = self.radii[i] + self.radii[j]
5     d        = (delta_x @ delta_v)**2 - (delta_v @ delta_v) * ((delta_x @
        delta_x) - R_ij**2)
```

Listing 1: Loop over all particles.

```
1 mask = np.arange(self.N-1)
2 mask[i:] += 1
3
4 r_ij = self.radii[mask] + self.radii[i]
5
6 delta_x = self.particles[:2,mask] - np.reshape(self.particles[:2,i],(2,1))
7 delta_v = self.particles[2:,mask] - np.reshape(self.particles[2:,i],(2,1))
8
9 vv = np.einsum('ij,ij->j',delta_v,delta_v)
10 vx = np.einsum('ij,ij->j',delta_v,delta_x)
11 xx = np.einsum('ij,ij->j',delta_x,delta_x)
12
13 d = vx ** 2 - vv * ( xx - r_ij**2 )
```

Listing 2: Vectorized calculation.

The calculation done in the simulation involves a bit more than this, but these blocks of code represents essentially all the calculation which is needed to determine if and when particle i collides with any other particle. Putting these two calculations into two functions `loop()` and `vect()` and comparing the time spent shows that the latter is approximately 100 times faster than the former, when doing a test on 50 000 particles. It should be noted however that the latter also requires a separate loop for pushing the new events into the queue, whereas the former does not. The speed-up is probably somewhat smaller than shown here.

```
1 >>> %timeit loop()
2 52.4 ms pm 260 us per loop (mean pm std. dev. of 7 runs, 10 loops each)

1 >>> %timeit vect()
2 429 us pm 23.3 us per loop (mean pm std. dev. of 7 runs, 1000 loops
  each)
```

3 Results and discussion

In this section, I will present the results from the simulations in the problems.

3.1 Speed distribution at equilibrium

The first test is devoted to investigating the distribution of the velocities of the particles after sufficiently long time. From statistical mechanics, we know that the velocities of a 2-dimensional gas at equilibrium will distribute according to Maxwell-Boltzmann's velocity distribution:

$$p(v) = \frac{mv}{k_B T} \exp\left(-\frac{mv^2}{2k_B T}\right), \quad (1)$$

where k_B is Boltzmann's constant, T the absolute temperature, and m the mass of the particles. It is evident that the particles will have to have the same mass for comparing with the known distribution in equation 1. The case of dissimilar masses is considered in section 3.2 and 3.3.

We initialise an system of particles with velocities $\mathbf{v} = [v_0 \cos \theta, v_0 \sin \theta]$ with $\theta \sim \mathcal{U}[0, 2\pi]$. The initial distribution of the velocities is therefore $\delta(v - v_0)$. After simulating the system until equilibrium is reached, i.e. the the average number of particle collision is $\gg 1$, the distribution is as shown in figure 1. For this simulation, the stop criterion used was that the average number of collisions surpassed 50. To get a larger number of independent samples, I have used the speeds at 20 different times in this simulation sampled such that all particles have collided a few times between each sample.




Figure 1: Distribution of velocities in a gas of 50000 particles.

To have a more quantitative measure of how good the fit is, divide the real line into k intervals consider the normalized ℓ^1 -norm of the difference at each interval

$$\text{err} = \frac{1}{2} \sum_{j=1}^k |d_j - p_j|, \quad (2)$$

where d_j denotes the density of samples in interval I_j and

$$p_j = \int_{I_j} p(v) dv$$

is the probability that an observation falls into interval I_j . It is evident from equation 2 that $\text{err} \in [0, 1]$. For the distribution shown in figure 1 $\text{err} \approx 0.0078527$.

3.2 Mixture of two gases

We repeat the exact same procedure as in section 3.1, except that we make half of the masses 4 times as big as the rest. The velocity distributions after equilibrium is reached is plotted in 2.




Figure 2: Distribution of velocities in a gas of 50000 particles. The red distributions correspond to the heavy particles, while the blue correspond to the lighter particles.

The average speed and kinetic energy is shown in table 1. These results suggests that although the masses are different, the mixture of the gases eventually reach equilibrium. The fact that the two gases reaches equilibrium is however most easily seen by actually plotting the time evolution of the average kinetic energy. This is shown in section 3.3, where we also simulate situations where this is *not* the case.

Table 1: Average speed and kinetic energy for the light and heavy particles

Mass	Average speed	Average kinetic energy
m	1.400236	12.466584
$4m$	0.702230	12.533416

3.3 Mixing in the presence of energy dissipation

To investigate, at least visually, whether the gas mixture reaches equilibrium we plot the time evolution of the averages of the energy of each species. In addition, we include two similar cases with the restitution parameter ξ set to 0.9 and 0.8. This is shown in figure 3. By the very definition of thermal equilibrium, e.g. the one given in [2, p. 3]

Thermodynamic equilibrium prevails when the thermodynamic state of the system does not change with time.

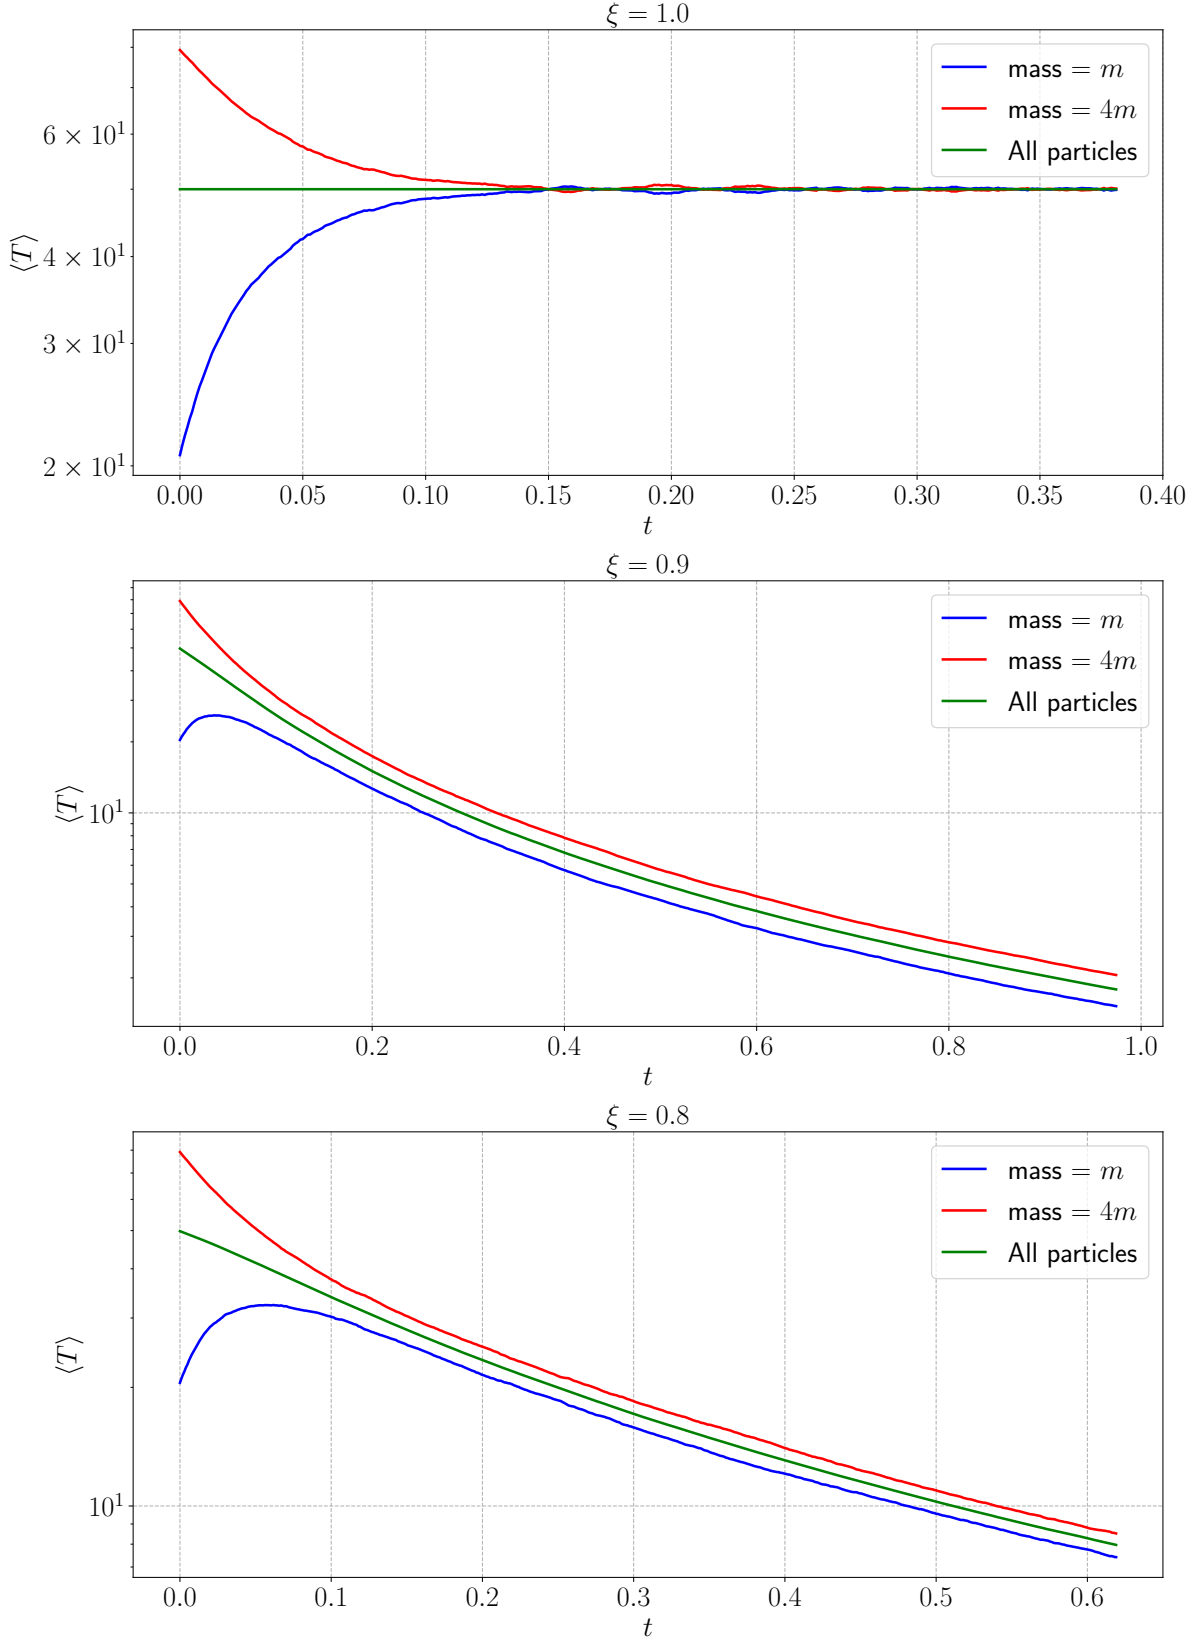
it is clear that the system cannot reach equilibrium in the presence of energy dissipation.

The plot in figure 3 shows that the two different particle species reach equilibrium only in the case where the energy is conserved. Furthermore, it is seen that the heavy gas molecules are always at higher temperature in the two latter cases. Although this fact is easily realised from elementary statistical mechanics, demonstrating it through these simulations is a good check that the correct physics is contained within the model made.

3.4 Crater formation from projectile impact

References

- [1] Aleksandar Donev, Salvatore Torquato, and Frank H. Stillinger. Neighbor list collision-driven molecular dynamics simulation for nonspherical hard particles. i. algorithmic details. *Journal of Computational Physics*, 202(2):737 – 764, 2005.
- [2] Kerson Huang. *Statistical Mechanics*. John Wiley & Sons, 2 edition, 1987.

Figure 3: Time evolution of energy averages for restitution parameter $\xi = 1.0, 0.9, 0.8$.